

To address this assignment, we'll break down the solution into two main sections: **implementation steps** and **explanation of algorithms (Linear Regression and Random Forest Regression)**. Additionally, we'll answer the specific questions about these algorithms at the end.

Let's start with the **Uber Fare Prediction project** implementation steps.

Project Implementation Steps

1. Data Pre-processing

- **Loading the Dataset**: Begin by loading the dataset using the `pd.read_csv()` function from Pandas. Here, you'll specify the path to the Uber dataset.

```
```python
import pandas as pd

data = pd.read_csv('uber_fare.csv')
```
```

- **Exploring the Dataset**: Inspect the data for null values, missing entries, and outliers using methods like `data.info()`, `data.describe()`, and `data.isnull().sum()`.

- **Handling Missing Values**: Address any missing data. Common techniques include removing rows with missing values or imputing missing values with the column mean or median.

```
```python
data.dropna(inplace=True) # Drop rows with null values
```
```

- **Feature Engineering**: Extract additional relevant features from the existing columns. For example, extract the day, month, hour, and day of the week from the `'pickup_datetime'` column to enhance model prediction.

```
```python
data['pickup_datetime'] = pd.to_datetime(data['pickup_datetime'])

data['hour'] = data['pickup_datetime'].dt.hour

data['day_of_week'] = data['pickup_datetime'].dt.dayofweek
```
```

```
'''
```

2. ****Identifying Outliers****

- Outliers can be identified using visualizations or statistical methods. For instance, using box plots for each feature (`fare_amount`, `passenger_count`, etc.) can reveal values outside the normal range.
- For continuous variables like `fare_amount`, you can use z-score or IQR (Interquartile Range) methods to detect outliers.

```
```python
from scipy import stats

data = data[(np.abs(stats.zscore(data['fare_amount'])) < 3)]
'''
```

## 3. **\*\*Check Correlation\*\***

- Generate a correlation matrix to observe relationships between variables. High correlation with the target (`fare\_amount`) and low correlation among features themselves are desirable.
- Visualize using `seaborn`'s `heatmap` to identify patterns.

```
```python
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10,8))

sns.heatmap(data.corr(), annot=True, cmap="coolwarm")
'''
```

4. ****Splitting Data into Train and Test Sets****

- Split the data into independent variables `X` (features) and the dependent variable `y` (target, `fare_amount`). Use `train_test_split` to divide the data into training and test sets.

```
```python
from sklearn.model_selection import train_test_split

X = data.drop(columns=['fare_amount', 'pickup_datetime'])
y = data['fare_amount']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
'''
```

#### 5. **\*\*Feature Scaling\*\***

- Use Standardization or Normalization to scale features to a uniform range, which can improve model performance. Use `StandardScaler` for continuous features.

```
'''python
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)
'''
```

#### 6. **\*\*Implementing Linear Regression\*\***

- Train a Linear Regression model using `LinearRegression` from `sklearn.linear\_model`.

```
'''python
from sklearn.linear_model import LinearRegression

lr_model = LinearRegression()

lr_model.fit(X_train, y_train)

lr_predictions = lr_model.predict(X_test)
'''
```

#### 7. **\*\*Implementing Random Forest Regression\*\***

- Train a Random Forest model using `RandomForestRegressor` from `sklearn.ensemble`.

```
'''python
from sklearn.ensemble import RandomForestRegressor

rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

rf_model.fit(X_train, y_train)

rf_predictions = rf_model.predict(X_test)
```

'''

## 8. **Model Evaluation**

- Evaluate both models using performance metrics like **R<sup>2</sup> score** and **Root Mean Squared Error (RMSE)**.

```
```python
from sklearn.metrics import mean_squared_error, r2_score

lr_rmse = mean_squared_error(y_test, lr_predictions, squared=False)

lr_r2 = r2_score(y_test, lr_predictions)

rf_rmse = mean_squared_error(y_test, rf_predictions, squared=False)

rf_r2 = r2_score(y_test, rf_predictions)
'''
```

Explanation of Algorithms

1. Linear Regression

Linear regression is a predictive modeling technique for predicting continuous outcomes based on input variables.

- **Simple Linear Regression**: Uses one independent variable to predict a dependent variable.

- **Multiple Linear Regression**: Uses multiple independent variables.

Linear Regression is represented by:

$$\hat{Y} = b_0 + b_1X + e$$

Where:

- \hat{Y} is the predicted outcome,

- b_0 is the intercept,
- b_1 is the coefficient for the independent variable X ,
- e is the error term.

Linear Regression uses **Least Squares** to minimize error by optimizing b_0 and b_1 values so that the difference between predicted and actual values is minimized.

2. Random Forest Regression

Random Forest Regression is an ensemble learning method that uses multiple Decision Trees to make predictions. Each tree in the Random Forest gives a prediction, and the final prediction is the average of all tree predictions.

How Random Forest Works:

1. **Bootstrap Sampling**: Multiple samples are drawn from the dataset with replacement.
2. **Feature Randomness**: Random subsets of features are used for each tree, reducing correlation among trees.
3. **Aggregation**: For regression, the mean of all predictions from the trees is used.

Advantages:

- **Accuracy**: Random Forests generally produce accurate results with low variance.
- **Overfitting Prevention**: By averaging multiple trees, Random Forest prevents overfitting compared to individual Decision Trees.
- **Feature Importance**: Random Forest can determine the importance of features, enhancing model interpretability.

Answers to the Questions

1. **Applications of Linear Regression:**

- Predicting house prices based on square footage and location.

- Estimating sales revenue based on marketing spend.
- Forecasting stock prices based on historical data trends.

2. **Important Functions in Linear Regression**:

- **fit()**: Trains the model on the training data.
- **predict()**: Generates predictions based on the fitted model.
- **score()**: Returns the R^2 score, indicating how well the model explains the variance in the data.
- **mean_squared_error()**: Computes the mean squared error, a measure of prediction error.

3. **How Random Forest Works and Advantages**:

- **Working**: Random Forests generate multiple Decision Trees, each trained on a different dataset sample and random subset of features. The predictions from these trees are averaged for regression tasks.
- **Advantages**:
 - **Handles High-Dimensional Data**: Performs well with a large number of features and instances.
 - **Reduces Overfitting**: Multiple trees average out individual variances, enhancing generalizability.
 - **Robust to Noise**: Random Forest is less sensitive to noise and outliers compared to other algorithms.

This explanation should provide a detailed understanding of each step and the algorithms used for Uber fare prediction. Let me know if you need further clarification on any part!

.....code explanation.....

Certainly! Let's dive into a detailed breakdown of each concept in the code and the approach for Uber fare prediction.

Step-by-Step Explanation

1. **Loading and Exploring the Dataset**

```
```python
```

```
data = pd.read_csv("Uber1.csv")
```

```
```
```

- **Loading the Data:** We use `pd.read_csv` to load the CSV file into a DataFrame, which allows us to manipulate and analyze the data. Each column in the DataFrame corresponds to a feature in our dataset:

- **Unnamed:** An index column, usually not useful for analysis.
- **key:** Unique identifier for each trip.
- **fare_amount:** The target variable, representing the fare for each trip.
- **pickup_datetime:** Date and time of the trip.
- **pickup_longitude and pickup_latitude:** Coordinates of the trip's starting location.
- **dropoff_longitude and dropoff_latitude:** Coordinates of the destination location.
- **passenger_count:** Number of passengers.

2. **Pre-Processing the Dataset**

Parsing Dates

```
```python
```

```
data["pickup_datetime"] = pd.to_datetime(data["pickup_datetime"])
```

...

- **\*\*Parsing Dates:\*\*** Converts ``pickup_datetime`` from a string to a ``datetime`` object, which is crucial for working with dates and times effectively (e.g., extracting day, month, or hour). This transformation simplifies any analysis based on time-related patterns.

#### #### Handling Missing Values

```
```python
missing_values = data.isnull().sum()

print("Missing values in the dataset:")

print(missing_values)

data.dropna(inplace=True)
```
```

- **\*\*Missing Values Check:\*\*** ``data.isnull().sum()`` returns a count of missing values for each column.

- **\*\*Drop Missing Values:\*\*** Since there are few missing entries, we can drop rows with ``NaN`` values using ``data.dropna(inplace=True)``. Alternatively, one could replace missing values with column means or medians (``data.fillna(data.mean())``).

---

### ### 3. **\*\*Identifying and Removing Outliers\*\***

Outliers can distort model training, especially in regression, where they can skew the mean and increase error.

#### #### Detecting Outliers Using Boxplots

```
```python
sns.boxplot(x=data["fare_amount"])

plt.show()
```
```

- **\*\*Boxplot for Visualization:\*\*** A boxplot shows the distribution of ``fare_amount`` and highlights any outliers as points outside the "whiskers" (which represent data within 1.5 times the IQR).



#### Calculating IQR (Interquartile Range)

```
```python
Q1 = data["fare_amount"].quantile(0.25)
Q3 = data["fare_amount"].quantile(0.75)
IQR = Q3 - Q1
threshold = 1.5
lower_bound = Q1 - threshold * IQR
upper_bound = Q3 + threshold * IQR
data_no_outliers = data[(data["fare_amount"] >= lower_bound) & (data["fare_amount"] <=
upper_bound)]
```
```

- **IQR Calculation:** The Interquartile Range (IQR) is used to identify outliers.
- **Q1 and Q3** represent the 25th and 75th percentiles of `fare\_amount`.
- **Threshold:** Values outside 1.5 times the IQR from Q1 and Q3 are considered outliers. This is a common threshold for outlier detection.
- **Filtering Outliers:** Only rows with `fare\_amount` between `lower\_bound` and `upper\_bound` are retained in `data\_no\_outliers`.

```

```

### 4. **Exploring Correlations Between Variables**

Correlations help understand relationships between variables, which can inform model selection and feature engineering.

```
```python
correlation_matrix = data.corr()
sns.heatmap(correlation_matrix, annot=True)
plt.show()
```

...

- **Correlation Matrix:** `data.corr()` calculates the Pearson correlation coefficient between numeric columns. Values range from -1 to 1, where:

- **+1** indicates a perfect positive correlation.

- **0** indicates no linear correlation.

- **-1** indicates a perfect negative correlation.

- **Heatmap Visualization:** A heatmap shows these correlations, helping identify which features might be strongly or weakly related to `fare_amount`. Higher correlations suggest more predictive power for that variable.

5. Preparing Data for Modeling

Defining Features (X) and Target Variable (y)

```
python
```

```
X = data[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude',  
'passenger_count']]
```

```
y = data['fare_amount']
```

...

- **Features (X):** We select relevant variables (`pickup_longitude`, `pickup_latitude`, `dropoff_longitude`, `dropoff_latitude`, `passenger_count`) to predict the fare.

- **Target (y):** `fare_amount` is the variable we want to predict.

Splitting Data into Training and Testing Sets

```
python
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

...

- **Train-Test Split:** Divides data into a training set (80%) and a test set (20%) to evaluate model performance. `random_state=42` ensures consistent results.

6. ****Building and Training Models****

Linear Regression

```
```python
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```
```

- ****Linear Regression:**** A basic algorithm that assumes a linear relationship between predictors and the target variable. It finds the best-fitting line by minimizing the squared differences between observed and predicted values.

Random Forest Regression

```
```python
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
```
```

- ****Random Forest Regression:**** An ensemble of decision trees, where each tree is trained on a random subset of data. Predictions are averaged to improve accuracy and reduce overfitting, making it more flexible for non-linear data than linear regression.

7. ****Evaluating Model Performance****

After training, we evaluate each model using metrics like R^2 (R-squared) and RMSE (Root Mean Squared Error).

Linear Regression Evaluation

```
```python
```

```
y_pred_lr = lr_model.predict(X_test)
```

```
r2_lr = r2_score(y_test, y_pred_lr)
```

```
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))
```

```
```
```

- **R² Score:** Indicates how much variance in the target variable is explained by the model. A higher R² (closer to 1) indicates a better fit.

- **RMSE:** Measures the average error in predictions. Lower RMSE means better performance.

Random Forest Regression Evaluation

```
```python
```

```
y_pred_rf = rf_model.predict(X_test)
```

```
r2_rf = r2_score(y_test, y_pred_rf)
```

```
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
```

```
```
```

- Same metrics (R² and RMSE) are calculated for the Random Forest model.

Comparing Models

```
```python
```

```
print("Linear Regression - R2:", r2_lr)
```

```
print("Linear Regression - RMSE:", rmse_lr)
```

```
print("Random Forest Regression R2:", r2_rf)
```

```
print("Random Forest Regression RMSE:", rmse_rf)
```

```
```
```

- **Interpretation:** The Random Forest model achieves a higher R² and lower RMSE compared to Linear Regression, indicating it captures more complex patterns and provides more accurate predictions.

```
---
```

Final Analysis

- **Why Random Forest Performed Better:** Uber fare data is influenced by complex, non-linear relationships (e.g., location coordinates), which are difficult for a linear model to capture. Random Forest's ensemble approach handles non-linearity and reduces overfitting by averaging predictions from multiple decision trees, providing more robust results.